

Operational AI starts one layer **too high.**

Most projects start with the agent.

That is the mistake.

5 tests before you deploy

The demo always looks convincing.

Ask why 20 orders are at risk.

"Likely supplier delays."

Confident. Impossible to prove.
Different tomorrow.

The trouble starts the
moment the answer has to
be trusted.

TEST 1 • CONTEXT

Does it know how your operation connects?

An order ties to materials, steps, suppliers, capacity, a promise date.

Or does it rebuild that from scratch every question?

TEST 2 • INSPECTABLE

Can someone open the rule behind "at risk"?

If the logic hides in a prompt, no one can challenge it, and no one can hand it over.

TEST 3 • EVIDENCE

Can every claim be traced to the records?

An answer you cannot trace is an answer you cannot defend.

TEST 4 • MEMORY

Will a validated insight still work next month?

Your business is constantly evolving.

Your AI needs evolve with it.

TEST 5 • CONTROL

Can it run inside your own data and controls?

Production data should not have to leave your environment to get an answer.

Build the model first. Then the **agent**.

Konekti models how your operation runs.

K-AI reasons over it to answer, explain causes, and turn validated insight into monitoring.

The goal is not a better chatbot. It's operational AI you can make decisions with.

Does your organization have AI that passes all 5 tests?

Most fail at least one.

You usually find out after you've already acted on a wrong answer.

Book a discovery call, we'll show you how you can deploy operational AI.

Follow Konekti for a closer look at each test.