June 10 2026

From insight to action

What Konekti does for ROSEN's service-delivery process — proven, and what it is worth

Prepared for [redacted], [redacted], and [redacted]. This note re-anchors what Konekti is for ROSEN after our three working sessions, separates what is already solved from what is not, shows four delivered customer results that map directly onto ROSEN's situation, and lays out the economics so the value can be weighed against the cost.

The one-line version

ROSEN has already won the hard fight — the data is extracted, modelled, and live. The value Konekti adds now is **speed, breadth, and reach:** letting process owners model and interrogate their own processes without queuing behind one person, evaluating far more of the analysis space than a human can by hand, on an object-centric foundation that answers the next question without rebuilding anything.

WHERE WE ACTUALLY ARE

Three sessions in: what's solved, and what isn't

ROSEN's process-mining capability is real and running. [redacted] built connectors into 10+ custom and legacy systems, wrote the extraction and transformation logic, derived a case ID that stitches a project across all of them, and stands up a new source connector in roughly two hours. Insights reach process experts today through a reporting dashboard. None of that is in question, and none of it is something Konekti would replace.

So the honest question from the 8 June session stands: if every step of the pipeline already works, where does Konekti help? The answer is not “we do what you do, but better.” It is that three specific constraints are quietly capping what the capability can become — and none of them is an extraction problem.

Constraint 1 — The capability runs through one person

Building and adapting a process model is [redacted]'s job today. There are around 30 process owners; roughly 5 are in scope for the first wave. Every new process and every model change lands in one queue. As [redacted] put it, if the process experts could model their own processes, it would *“solve this capacity issue.”* That is the constraint in his own words.

Constraint 2 — One object means one focus for questions

Today ROSEN's model follows a single object — the report, tracked by one case ID across systems. That is a **case-centric** view, and it answers the report-cycle-time question well. But the moment the business asks something else — throughput for a specific ILI tool, rework traced to a sensor fault, a work group's variance — the model has to be rebuilt around a new case notion. Object-centric process mining models tools, reports, customers, runs, and people, with relationships intact, and projects any process from that data model.

Constraint 3 — Insight is capped by what a human thinks to look at

ROSEN's analysis surfaces what process owners already suspect; the questions are shaped by who is in the room. Konekti's agent evaluates the analysis space directly — it tests hundreds of paths and surfaces the ones that matter, including the ones nobody thought to ask about. That is the difference between confirming a hunch and finding the thing you didn't know to look for.

THE THING THAT'S EASY TO GET BACKWARDS

Two things are true about object-centric

It would be easy to read the last few pages as “go object-centric and everything gets easier.” Downstream, that is true — but it skips the catch, and ROSEN already knows the catch, so it is worth stating plainly. There are two distinct layers here, and the second is the one that changes the decision.

Layer 1 — Konekti builds an object-centric model natively

The event store is object-centric at its core, not case-centric data with an object-centric label on top. Tools, reports, runs, customers, and people are modelled as distinct objects with their real relationships intact — not one case with everything else flattened onto it. That native structure is what opens up reuse, cross-object root cause, and the breadth the agent reasons over. Everything in the four examples below stands on this.

Layer 2 — Konekti makes building that model the easy path, not the hard one

Here is the part that matters most for ROSEN: building an object-centric model by hand is **harder** than building a case-centric one. More objects, more relationships, more to get right. That difficulty is exactly why most teams stop at case-centric — it is the pragmatic choice, and it was the right first move when proving the capability solo, as ROSEN did. The catch is that the easier thing to build is also the thing that can only ever answer questions about one focus.

> **Konekti removes the trade-off.** It makes the harder, more powerful model the *easy* one to build — no-code, visual, agent-assisted. ROSEN gets the object-centric foundation that unlocks everything downstream, without paying the object-centric build cost that normally forces teams to settle for case-centric. The hard part becomes the easy part; the easy part was never the point.

PROVEN RESULTS

Four examples of things Konekti has already done

These are delivered customer outcomes, each chosen because it maps onto a part of ROSEN's situation: a fully custom manufacturing estate, conversational analysis at speed, adoption that reached leadership, and self-service modelling that freed the bottleneck engineer.

BMW — Root cause across a fully custom manufacturing system

BMW ran production on a heavily customised manufacturing estate — the kind of environment, like ROSEN's, where no off-the-shelf process-mining connector fits and event data has to be modelled from systems never built to be mined. The problem was not a missing dashboard. It was that rework, bottlenecks, and quality-assurance issues kept recurring and no one could pin down why, because the cause crossed systems and stations that each held only part of the story.

Konekti modelled the production reality object-centrically — orders, components, stations, quality checks, and the relationships between them — then let the team decompose the problem across that connected structure. **Multiple distinct root causes surfaced behind what had looked like one vague "rework problem,"** each traceable to specific stations, handoffs, and conditions in the evidence. The team moved from a recurring symptom to a ranked set of causes they could act on, and drove targeted improvement against each.

3x less effort	100%	15min
to transform data (model, validate, code)	custom systems — no standard ERP connector	to onboard process owners

Relevance to ROSEN: this is the closest analog to your estate. Fully custom, non-process-aware systems are the documented hard case for process mining; most tooling assumes SAP. BMW shows Konekti's approach holds up exactly where custom code and standard tools don't.

ABB — Conversational analytics at a speed and breadth a person can't match

A global process analyst at ABB needed to work out where their inventory management process was losing time and value. That kind of analysis — testing hypotheses, slicing by segment, chasing each lead, cross-referencing — typically takes a skilled analyst around four weeks.

In a single prompt, Konekti's agent **evaluated and prioritised more than 113 distinct analysis paths and surfaced the top 15 for deeper agentic analysis and human review.** The analyst completed in four hours the analysis that normally takes four weeks — not by automating a known report, but by searching a breadth of the analysis space no individual could cover by hand, and handing back a ranked shortlist grounded in the evidence.

113	Top 15	4 weeks → 4 hrs
analysis paths evaluated in one prompt	surfaced for deeper review	same analysis, ~40× faster

Relevance to ROSEN: this is the agentic approach we discussed, doing the thing your current approach can't — not answering one question faster, but covering a search space wide enough to surface what process owners are blind to. On a case-centric event log an agent can only reason inside one case; this breadth requires the object-centric store underneath.

Industrial lighting manufacturer — Adoption that reached the C-suite

A global industrial lighting manufacturer wanted leadership — not just analysts — to understand how work actually flowed through the business. This is the adoption problem in its hardest form: getting senior decision-makers to engage directly with process intelligence rather than receiving a filtered slide three weeks later.

The reusable object-centric model was built by **an analyst who had never looked at the data before, with no input from process owners, in under 10 hours.** From it, C-suite leadership surfaced **five completely separate strategic improvement pathways, on their own** — five distinct directions for the business, from one model, read directly by the people who own the decisions. Because the model is object-centric, it is reusable and expandable: the organisation keeps adding data sources and extends the same model over time rather than starting over.

<10 hrs	0	5
to a reusable model, analyst new to the data	process-owner input required to build it	strategic pathways surfaced for leadership

Relevance to ROSEN: adoption is the wall the 8 June session kept hitting — insights exist, but action stalls. The route through is reach: when leadership and process owners interrogate the model themselves, the insight that drives action no longer depends on one person's availability. And the model an analyst builds in a day keeps compounding instead of being thrown away. Shifting analysis to process owners empowers them with trust and understanding of the insights and accountability for the resulting improvement actions.

Takeda — Lowering the barrier so analysts self-serve

At Takeda, deviation management was modelled the way serious process work used to require: a data engineer built the model in code, and a process analyst worked alongside them — the analyst even had to become passable at reading SQL just to follow the logic. Every new question queued behind the engineer. That is precisely ROSEN's Constraint 1.

With Konekti, the data modelling moved off the engineer entirely. Process analysts built and adapted the model themselves in Konekti, no code — they had the hypotheses, so they built the objects, processes, and rules directly, validating each step. The backlog cleared while process mining demand grew. And critically, the engineer's freed time went to higher-value work: **more complex modelling, integrating new data sources, and building the automations that supported the action steps coming out of the insights** — closing the loop from finding a problem to acting on it.

Analysts	1 engineer	Insight → action
now build models themselves, no SQL	freed for complex work + automations	loop closed with built automations

Relevance to ROSEN: Constraint 1, solved, with a named pharmaceutical manufacturer. The bottleneck moves off [redacted]; process owners self-serve; the engineer's time redirects to landing new data and building the automations that make insights get acted on — the exact thing the 8 June session identified as missing.

THE MAPPING

How each result lands on ROSEN

The four examples are not a general capability brochure. Each answers a specific thing ROSEN said in session.

ROSEN reality (from session)	Proven by	What it means for ROSEN
Fully custom, non-process-aware systems; no off-the-shelf connector fits	**BMW**	Konekti's object-centric approach holds up in exactly the estate that defeats manual code and standard tools.
"We only look at one object — we'd have to do it again, differently"	**ABB**	One model, breadth no human matches; the next question is a projection, not a rebuild.
Insights exist but action stalls; only [redacted] can model	**Lighting mfr**	Reach to leadership and a model built in a day that compounds — adoption stops depending on one person.
Process owners can't self-serve; engineer is the bottleneck	**Takeda**	Modelling moves to analysts; engineer's time goes to new data and action-step automations.

> **Why this matters for the agent ROSEN wants**
> [redacted]'s concern — that ROSEN may host its own agents in 12 months — is fair. But an agent is only as good as the structure beneath it. The ABB result (100+ paths in one prompt) is not a property of the language model; it is a property of reasoning over an object-centric store in process vocabulary. An agent on a case-centric event log can only reason inside one case — it hits the same wall the model does. A home-built agent on today's structure inherits today's ceiling. The differentiator was never "who has an agent"; it is what the agent stands on.

THE RE-ANCHOR

What Konekti is, in ROSEN's terms

Strip away the category language. For ROSEN specifically, Konekti is three things sitting on top of the data you already extract:

Layer	What it is for ROSEN	Proven by
Native object-centric model	Event store that is object-centric at its core — tools, reports, runs, customers, people modelled once with relationships intact, projected into any process on demand.	BMW, Lighting mfr
Easy to build, by design	The harder model made the easy one: no-code, visual, agent-assisted. Object-centric without the object-centric build cost — and process owners do it, not just engineers.	Takeda
Process-native agent	Ask process questions in plain language; breadth no human matches, answers grounded in the model. Runs in your cloud, your tokens.	ABB

Deployment is unchanged from what was discussed: Docker in ROSEN's **private cloud**, execution as SQL/Spark in your own data platform, no operational data leaving your perimeter, your own LLM tokens. For a public-safety-sector organisation, that on-premise posture is the price of entry — and it is met.

What Konekti is explicitly NOT doing here

- **Not touching extraction.** Your connectors into legacy systems stay yours.
- **Not replacing the dashboard wholesale.** It changes who can build the view behind it and how fast it adapts.
- **Not selling an agent as a standalone gadget.** The agent is the payoff of the reusable model, as ABB shows.
- **Not a multi-year transformation.** The entry is one scoped use case — the report process you already chose.

WHAT COMES WITH IT

A standing line to people who've done this for years

Konekti's operating model is not software-and-goodbye. The engagement includes a daily 15-minute check-in, hands-on onboarding workshops, and user training. That time is ROSEN's to use however is most useful in the moment — unblocking a modelling decision, pressure-testing an approach before committing to it, or working a hard edge case through live.

In practice it is continuous access to a team that has built object-centric process intelligence across many industries and has already hit the pitfalls before you reach them. ROSEN has built a serious capability from the ground up, fast, with a small team. This puts experienced people alongside that work — not to take it over, but to give it leverage: what is becoming best practice across the field, what breaks at scale, what the next object-centric question usually turns out to be. **Fewer dead ends, less rework, and a faster path to getting it right the first time.**

> **Why it is worth naming:** the hardest part of moving from a solo-built pipeline to a broader capability is rarely the tooling — it is not knowing which of today's decisions will cost you in eighteen months. A standing line to people who have already made and corrected those decisions is how that risk gets retired. It is the difference between learning the pitfalls by hitting them and learning them before you do.

THE ECONOMICS — FOR THE P&L CONVERSATION

Tying this to ROSEN's bottom line

In the June 8 session, the ROI question kept landing as “not my position to answer.” That is reasonable — it is [redacted]'s question, because [redacted] owns the service view. Following a call with [redacted], what follows is not an ROI claim; it is the **model [redacted] can put numbers into,** using figures only ROSEN holds — most of which his team has already calculated.

One mechanism, three pools of value

Lifting assessment-and-analysis capacity is a single lever — it shows up three ways: cost saved, penalties avoided, and revenue enabled. Each formula uses figures only ROSEN holds, most of which [redacted]'s team has already calculated.

Pool 1 — Penalty avoidance (seasonal, and likely the largest)

ILI runs peak in winter and summer. Off-peak, throughput is fine. On-peak, the analysis side becomes the bottleneck, on-time delivery slips, and contractual penalties hit. Relieve the bottleneck and most of those penalties never occur.

> **Avoided penalty cost = breaches/year × avg cost/breach × 0.8**
>
> *0.8 = the share of breaches driven by the high-season analysis bottleneck — the avoidable portion. ROSEN sets the breach count and cost.*

Pool 2 — Efficiency & rework (steady-state, from ROSEN's own data)

Less rework and more efficient report writing means fewer man-hours per report. ROSEN has already quantified gain areas; Konekti's breadth-of-search (the ABB mechanism) surfaces more the team hasn't found yet. Nothing counts until it is implemented, so the figure is discounted to reflect that.

> **Annual saving = total hours /report × (major gain % × 0.6) × (PO gain % × 0.8) × loaded hourly rate × reports/year**
>
> *Illustrative shape:*
> *1. Say ROSEN has identified ~1 major gain area that saves 15% time, assume a realization of 0.6 as major changes are more difficult to implement.*
> *2. Say process owner’s agentic analysis surfaces ~5 gains that save 5% time, assume a realization of 0.8 due to easier implementation and improved action.*

Pool 3 — Capacity / value generation (the upside on top)

Analysis capacity is the governor on how many ILI runs ROSEN can confidently take on. Free up hours and they can grow run volume — but a bigger book re-hits the penalty ceiling at the new level, so the honest figure is net of the new penalties that higher volume brings.

> **Net new value = (value/ILI run × additional runs enabled) − new penalties at higher volume**

The framing for [redacted]

The efficiency saving (Pool 2) is real but modest, and ROSEN can compute it today from existing findings. The number that likely justifies this on its own is Pool 1 — penalties are contractual cash, they cluster in the peak seasons, and they are caused by precisely the bottleneck Konekti relieves. Pool 3 is the growth case: once analysis is no longer the constraint, the top of the funnel can expand. **Set against a client-led POV at six weeks with a 3-month Konekti subscription at €7,500/month, the penalty pool alone tends to settle the question.**

Alternatively, a Konekti-led serviced POV is also possible for a flat fee.

On the “we might build it ourselves” question

It deserves a straight answer. If ROSEN's infrastructure team will build and maintain object-centric, process-native modelling and a grounded agent in-house, and has it on the roadmap, that is a legitimate path. But the session surfaced that the agent work is not on the infrastructure team's roadmap, and process mining is not in their target range, and the timeline is 6–12 months for general agent access before any process-specific agentic work begins. The choice is not “buy now vs. build the same thing soon.” It is “move on the bottleneck this quarter vs. wait a year for capacity that isn't yet scoped.” That is a value-of-time question, and it is [redacted]'s to weigh.

To be brutally honest, process specific agents are likely 2 years out at a minimum, and that assumes an internal capacity to build the data infrastructure and in-house expertise in multi-agent orchestration to query and mine the data models.

PROPOSED NEXT STEP

A scoped POV that proves it on one process

Keep it small and tied to the process ROSEN already chose. The aim is not a transformation — it is one honest test on real data, proving the same three things the examples above proved elsewhere.

Element	Proposed scope
Use case	The existing report / service-delivery process — already mined and tied to revenue.
Prove self-service and re-use (Takeda and BMW)	Two process owners build/adapt a process themselves, no SQL, agent-assisted first drafts.
Prove reuse (Lighting)	From the same store, project one second view (e.g. tool-level or work-group) without rebuilding.
Prove breadth (ABB)	Run the agent across the process; measure analysis paths surfaced and time vs. the current manual approach.
Success metric	Self-serve modelling time vs. today's queue; second-view build time; analysis breadth and speed; one standardisation action actually shipped.
Commercial	Six-week POV, three months' access. On-prem, your tokens. Service or self-serve path per ROSEN's preference.

What we are asking for next: a 30-minute follow-up with [redacted], [redacted], and [redacted] together to pressure-test the economics with ROSEN's own margin-per-report figure, and to confirm the two process owners for the POV. If the throughput math doesn't clear the cost bar on your numbers, that is a clean no — and worth knowing fast.