From disconnected sources to revenue you can explain

A phased approach for Innova Market Insights

Prepared for Andrzej Murajew, Digital Marketing Strategist, Innova Market Insights
Prepared by Adam Broniewski, Customer Success Manager, Konekti — following our discovery call, 4 June 2026

Date: 10 June 2026
Valid until: 10 July 2026
Confidential

What we heard

You have the metrics on paper — CAC, LTV, LTV:CAC, cost per lead, cost per action, cost per sale — but you cannot produce them without people stitching numbers together in Excel by hand. Ten systems, and as you put it, *nothing connects to nothing*. GA4 does not talk to the website, the website does not talk to the dashboard, the dashboard does not talk to the CRM, and Dynamics does not talk to any of it.

The cost of that gap is not abstract. You recently lost one of your largest clients, with no warning and no way to look back and understand what happened beforehand. That is the problem worth solving first: **not knowing who is about to leave until they are already gone.**

> **The decision this should change**
>
> Today, when a client goes quiet, you find out at renewal. The goal of a first engagement is to surface the warning signs **while there is still time to act** — and to do it from data you already own, without waiting a year for all ten systems to be connected.

How we would approach it

We do not recommend a one-year, connect-everything program. Those projects spend twelve months building plumbing before anyone sees a single answer, and by then the question has often changed.

Instead we build the end-to-end picture **one frame at a time**. Each frame is a self-contained slice of your revenue process — a single question, the minimum data needed to answer it, and a working result you can act on. We model that slice, validate it against real records, and stand up conversational analysis on top, so you can ask follow-up questions in plain language. Then we add the next frame.

The payoff: your **first insight in roughly six weeks**, not after a year — while the next frame is already being built in parallel.

Every frame has the same four layers

This matters for your budget, because the layers are not all priced the same — and some of them your own team may choose to take on.

Layer	What happens	Who is best placed to do it
1 — Data integration	Connect the relevant source(s) into your Redshift warehouse: connectors, tables, refresh frequency, schema-change handling.	This is general data-engineering effort; where you have internal capacity, this is where you save the most cost.
2 — Modelling	Turn raw tables into a reusable model of your revenue process: leads, deals, customers, renewals, and how they connect.	Konekti. This is our core.

Layer	What happens	Who is best placed to do it
3 — Analytics & AI	Root-cause analysis, the metrics you listed, anomaly signals, and a natural-language Q&A agent that understands your process model.	Konekti. This is our core.
4 — Handover	Documentation, training, and an agreed support model so your team can run and extend it.	Konekti, lightweight.

The revenue process, in work packages

We have split your funnel into five packages along the path from first touch to renewal. Each one delivers a usable answer on its own, and each is scoped so it can be answered with the systems that package connects — nothing borrowed from a package that has not been built yet.

Two things drive the order. First, **the churn question comes first**, because that is your live pain. Second, **the easier-to-connect packages come before the harder ones**. A package that needs only the CRM can start almost immediately; one that needs every ad platform stitched back to the CRM cannot, however interesting its question. So the sequence below runs from least to most integration-heavy, with the attribution package (WP3) deliberately later than its funnel position would suggest — it is the one gated on the most data work.

These are our opening assumptions, not a fixed plan. Tell us where our hypotheses are wrong and we will resequence. If your gut says the earliest churn signal is people going quiet in the CRM, WP1 is already the right start; if you believe it is disengagement from campaigns, that signal lives in WP4's data and we would weigh bringing it forward against the heavier integration it carries.

WP1 Renewals & churn signal ★ START HERE	
Key question	**Which current clients are most likely to churn next — and why didn't we see the last one coming?**
Systems used	Dynamics 365 (CRM): renewal records, deal/opportunity records, logged account activity. The custom CRM-analysis dashboard's underlying data if it holds anything the CRM does not.
Integration weight	**Lightest. One system, already partly proven (it feeds your Power BI dashboard). Can start almost immediately.**
Why first	It is the live wound, and it is answerable from data you already own. Understanding who did not renew, and what their record looked like beforehand, is what builds the pattern for who is at risk now.
Working hypotheses	• Accounts that churned went quiet in the CRM first — fewer logged interactions, longer gaps since last contact, fewer open opportunities — in the months before renewal lapsed. • A renewal quote created late, or never created, is an early warning of a lost account.
Answerable here?	Yes — both hypotheses rely only on CRM data this package connects. (The product- and campaign-engagement signal — logins, report downloads, newsletter opens — is a real churn predictor too, but it lives in systems WP4 and WP5 connect, so it is added there as an upgrade.)
Minimum data	Dynamics renewal and deal records. Where getting a live connection is hard, a one-off export of active vs. lapsed clients (with renewal dates and last-activity dates) is enough to start.
Answer you get	A churn-risk view across the current book, the CRM signals that precede a loss, and the cases to call this quarter.

WP2 Lead → deal (acquisition)	
Key question	Which channels and sources actually produce closed deals — not just leads?
Systems used	Dynamics 365 (CRM): deal pipeline plus the lead-source field already captured on each lead/opportunity.
Integration weight	**Light. Same system as WP1; no new connectors, just more of the CRM modelled.**
Working hypotheses	• A small number of sources drive most won revenue, while budget is spread across all of them. • Cost-per-lead and cost-per-deal differ sharply by source once you follow the lead all the way to a close.
Answerable here?	Yes, as far as the CRM's own lead-source field allows. Tying deals back to the specific ad or campaign that created them is WP4's job — this package answers it at the source/channel level the CRM already records.
Minimum data	CRM deal pipeline plus lead source. First-pass cost-per-lead and cost-per-deal become possible here.
Answer you get	A first, defensible read on CAC and cost-per-deal by source — the start of the reallocation case (the 80/20 you described).

WP3 Deal → cash (revenue realisation)	
Key question	What is each client and channel actually worth once you count realised, paid revenue — not booked deal value?
Systems used	Dynamics 365 plus wherever realised revenue lives — order/contract/invoice records, which may be a separate Dynamics module or a finance system (to confirm).
Integration weight	**Medium. Adds one new data area (finance/orders) and the link back to CRM accounts, but not the full channel stack.**
Working hypotheses	• True LTV and LTV:CAC need realised revenue, linked to accounts through orders, contracts, and invoices — not booked deal value.
Answerable here?	Yes, once the revenue source is connected and joined to CRM accounts. Builds directly on the WP1–WP2 customer model.
Minimum data	Order / contract / invoice records and their link to CRM accounts.
Answer you get	Real LTV and LTV:CAC by client and source, closing the loop between acquisition cost and the revenue it ultimately produced.

WP4 Campaign → lead (marketing attribution)	
Key question	Which marketing activity brings in the leads that become customers?
Systems used	GA4, Search Console, the ad platforms (Google Ads, Bing, LinkedIn, Meta), Eloqua — plus identity stitching back to the CRM. Nearly the whole stack.
Integration weight	**Heaviest. Needs the most sources connected and the hardest problem solved — stitching click/visit identity through to the CRM. This is why it sits here, not earlier: its funnel position is upstream, but its data dependency is the deepest.**
Working hypotheses	• Engagement channels (newsletter, website, LinkedIn) differ greatly in the quality of leads they generate, not just the volume. • UTM / click-ID continuity from the ad platforms into Eloqua and Dynamics is currently broken or partial — a key thing to confirm (see technical questions).
Answerable here?	Yes, but only once the channel data is connected and identity is resolved. This is also where the product/campaign-engagement churn signal (logins, downloads, newsletter opens) can be fed back to sharpen WP1.
Minimum data	GA4, Search Console, ad platforms, Eloqua, plus identity stitching to the CRM.
Answer you get	Multi-touch attribution from first click to closed deal, and CPL / CPA / CPS across channels under first-click, last-click, and data-driven models.

WP5 Customer service & engagement (optional extension)	
Key question	Do support interactions and post-sale engagement predict churn earlier than commercial signals alone?
Systems used	Ticketing / customer-service system (to confirm one exists and is accessible), joined to the customer model from WP1.
Integration weight	**Optional add-on. One new system, joined to an existing model.**
Working hypotheses	• Support tickets, response times, and unresolved issues are leading indicators of churn that the commercial data misses.
Answerable here?	Yes, if a ticketing system exists and its data is reachable.
Minimum data	Ticketing / support system, joined to the customer model built in WP1.
Answer you get	A richer churn-risk model that combines commercial and service signals — a direct upgrade to WP1.

How the packages compound

WP1 gives you a churn view from CRM data alone. WP2 adds which sources bring the deals. WP3 weights it all by real, realised revenue. WP4 brings the full channel picture — multi-touch attribution, and the campaign-engagement signal fed back to sharpen the WP1 churn model. WP5 adds the service signal. You get a usable answer after the first package, and a sharper one after each that follows — instead of one big reveal at the end.

What we need before we can put numbers on this

We deliberately have not priced these packages yet. The effort in each one is driven almost entirely by Layer 1 — the state of your data — and only your side can tell us that. A package that looks like the obvious place to start can turn out to be the hardest to connect. Once we have the answers below, we can overlay effort and cost onto the priority order, and you can decide the trade-off: pay a premium for the hardest-but-highest-impact frame, take the cheapest, or balance the two.

Where you want us to focus — for you

These are about direction, not numbers. They tell us where to point the first package and how you would sequence the work.

1. What is the single decision you most need to make today and cannot? (If you could have one answer next month, what would it be?)
2. Where would you start? You know this business better than we do — which part of the funnel (renewals, acquisition, attribution) do you believe holds the biggest problem?
3. What is your hypothesis for the churn you have seen? When you lost accounts, what is the teams gut instinct on the early signal we would have caught — going quiet in the CRM, disengaging from campaigns, support issues, something else?
4. Across your marketing channels, is spend spread fairly evenly, or concentrated in a few? (This tells us how much value a reallocation insight could unlock — we just need the shape, not the amounts.)
5. How does your team define a 'won' client and realised revenue — is it signed, invoiced, or collected? (This decides what WP3 actually measures, and it is a business call, not a technical one.)
6. What is the target timing for a first usable result?
7. Is there a budget envelope we should design the first package within for year one? (We are not quoting yet — this just keeps our first proposal realistic rather than over- or under-built.)

On the ten systems — quick confirmation

Your note listed ten sources: GA4, Search Console, Google Ads, Bing Ads, LinkedIn, Facebook/Instagram, Twitter, Oracle Eloqua, Microsoft Dynamics 365, and an internal CRM-analysis dashboard. Before we size anything, one thing to confirm:

8. Are these ten genuinely separate sources, or are they sharing an origin? For example: is the internal CRM-analysis dashboard its own data store, or a view built on Dynamics? Are Facebook and Instagram sharing a single Meta connection? Is GA4 your only website analytics, or is there a separate web/Matomo source as well? Knowing which are truly distinct changes the integration count materially.

Technical — for your tech team

These size Layer 1 in each package. Most can be answered in a short call with whoever owns the warehouse.

9. Is Redshift live today, and which of the systems already have a working connection into it? In particular — the Dynamics → Power BI dashboard suggests a Dynamics connection exists: is that data already landing in Redshift, or only feeding Power BI, and at what refresh rate?
10. Are UTMs and click IDs passed from the ad platforms (Google Ads, Bing, LinkedIn, Meta) through into Eloqua and Dynamics? This is the single biggest driver of how hard the attribution package (WP4) will be.
11. How is a person or company identified across systems — is there a shared key, or is cross-system identity resolution unsolved today?
12. Is there any transformation layer on Redshift that we would need to build on or integrate with, or would we be setting that structure up?
13. Where is realised revenue physically stored — same Dynamics instance, a separate module, or a finance system? (The definition of what counts as realised revenue is the business question above; this one is only about where the data lives.)
14. Do you have a ticketing / customer-service system, and is its data accessible? (Relevant only to the optional WP5.)
15. Do you want us to operate the extraction pipelines on an ongoing basis, or only design them and hand them to your team to run?

Proposed next step

Pick the one or two packages that match your highest-priority question. We take the technical answers above for just those, return an effort-and-cost estimate for them, and book a session with your tech lead to confirm the data reality. That keeps the first commitment small, concrete, and tied to a decision you actually need to make.

Konekti turns fragmented operational data into grounded, inspectable decision intelligence. Your data stays in your environment — Konekti generates execution logic that runs inside your own data platform.

getkonekti.io · Eindhoven, The Netherlands